

Tamas Flesch Thesis - Decision Tree step

LMJU - UpGrad - DS

Fifa 23 Ultimate Team player price prediction based on the player's attributes

Table of contents

- Imports
- Read data
- Inspect the data
- Build decision tree model
- Using a random forest
- Grid Search experiment

Imports

```
In [2]: # Importing the required libraries
import pandas as pd, numpy as np
import matplotlib.pyplot as plt, seaborn as sns
%matplotlib inline
```

Read data

```
In [4]: # Reading the csv file and putting it into 'df' object.
df = pd.read_csv('futbin.csv')
```


Inspecting the data

```
In [5]: df.head()
```

Out[5]:

	Name	Rating	Price	Skills_Star	Weak_Foot_Star	Pace / Diving	Shooting / Handling	Passing / Kicking	Dribbling / Reflexes	De / S
0	Pelé	98	3270000.0	5	4	95.0	96	93	96	
1	Lionel Messi	98	4350000.0	4	4	93.0	98	97	99	
2	Lionel Messi	98	4640000.0	4	4	94.0	97	96	99	
3	Karim Benzema	97	1850000.0	4	5	92.0	97	90	94	
4	Kylian Mbappé	97	9750000.0	5	4	99.0	96	88	98	

5 rows × 105 columns




In [32]: `df_names = df.pop('Name')`

In [33]: `df.head()`

Out[33]:

	Rating	Price	Skills_Star	Weak_Foot_Star	Pace / Diving	Shooting / Handling	Passing / Kicking	Dribbling / Reflexes	Defense / Speed	Pl Pos
0	98	3270000.0	5	4	95.0	96	93	96	60	
1	98	4350000.0	4	4	93.0	98	97	99	40	
2	98	4640000.0	4	4	94.0	97	96	99	40	
3	97	1850000.0	4	5	92.0	97	90	94	45	
4	97	9750000.0	5	4	99.0	96	88	98	44	

5 rows × 104 columns



In [34]: `df_names.head()`

Out[34]:
0 Pelé
1 Lionel Messi
2 Lionel Messi
3 Karim Benzema
4 Kylian Mbappé
Name: Name, dtype: object

Build a decision tree model

In [20]: `from sklearn.tree import DecisionTreeRegressor`

In [21]: `dt = DecisionTreeRegressor(random_state=42, max_depth=4, min_samples_leaf=10)`

```
In [35]: np.random.seed(0)
df_train, df_test = train_test_split(df, train_size=0.8, random_state=100)
```

```
In [36]: df_train.shape, df_test.shape
```

```
Out[36]: ((5117, 104), (1280, 104))
```

```
In [37]: df_test.head()
```

```
Out[37]:
```

	Rating	Price	Skills_Star	Weak_Foot_Star	Pace / Diving	Shooting / Handling	Passing / Kicking	Dribbling / Reflexes	Defense / Speed	Ph Posi
2691	69	250.0	3	3	60.0	69	51	62	28	
3785	66	200.0	3	3	73.0	65	49	66	22	
4989	62	200.0	0	3	64.0	61	61	64	40	
1222	75	450.0	3	3	53.0	53	74	79	61	
3640	66	200.0	2	3	65.0	61	58	64	65	

5 rows × 104 columns

```
In [38]: scaler = MinMaxScaler()
```

```
In [39]: df_train['Price'] = scaler.fit_transform(df_train[['Price']])
df_test['Price'] = scaler.transform(df_test[['Price']])
```

```
In [40]: df_train.Price.describe()
```

```
Out[40]: count    5117.000000
mean         0.001729
std          0.022844
min          0.000000
25%          0.000000
50%          0.000010
75%          0.000053
max          1.000000
Name: Price, dtype: float64
```

```
In [41]: y_train = df_train.pop("Price")
X_train = df_train

y_test = df_test.pop("Price")
X_test = df_test
```

```
In [42]: X_test.shape, X_train.shape
```

```
Out[42]: ((1280, 103), (5117, 103))
```

Fit the DT model

```
In [43]: dt.fit(X_train, y_train)
```

```
Out[43]: DecisionTreeRegressor(max_depth=4, min_samples_leaf=10, random_state=42)
```

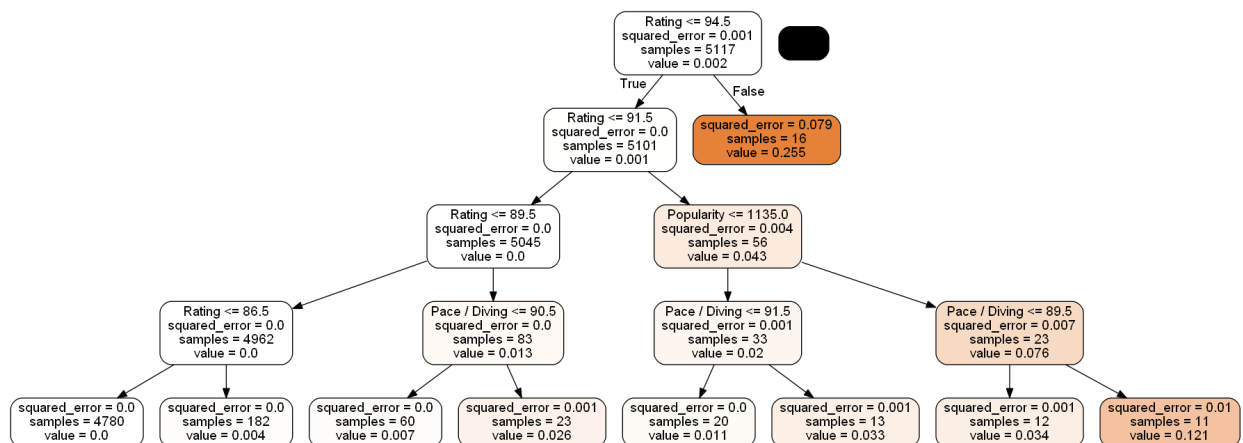
```
In [44]: from IPython.display import Image
from six import StringIO
from sklearn.tree import export_graphviz
import pydotplus, graphviz

dot_data = StringIO()

export_graphviz(dt, out_file=dot_data, filled=True, rounded=True,
                feature_names=X_train.columns)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

```
Out[44]:
```



```
In [45]: y_train_pred = dt.predict(X_train)
```

```
In [47]: from sklearn.metrics import r2_score
```

```
In [48]: r2_score(y_train, y_train_pred)
```

```
Out[48]: 0.46511633392209173
```

```
In [49]: y_test_pred = dt.predict(X_test)
```

```
In [50]: r2_score(y_test, y_test_pred)
```

```
Out[50]: 0.41381406125743336
```

Using Random Forest Regressor

```
In [51]: from sklearn.ensemble import RandomForestRegressor
```

```
In [52]: rf = RandomForestRegressor(random_state=42, n_jobs=-1, max_depth=5, min_samples_leaf=1)
```

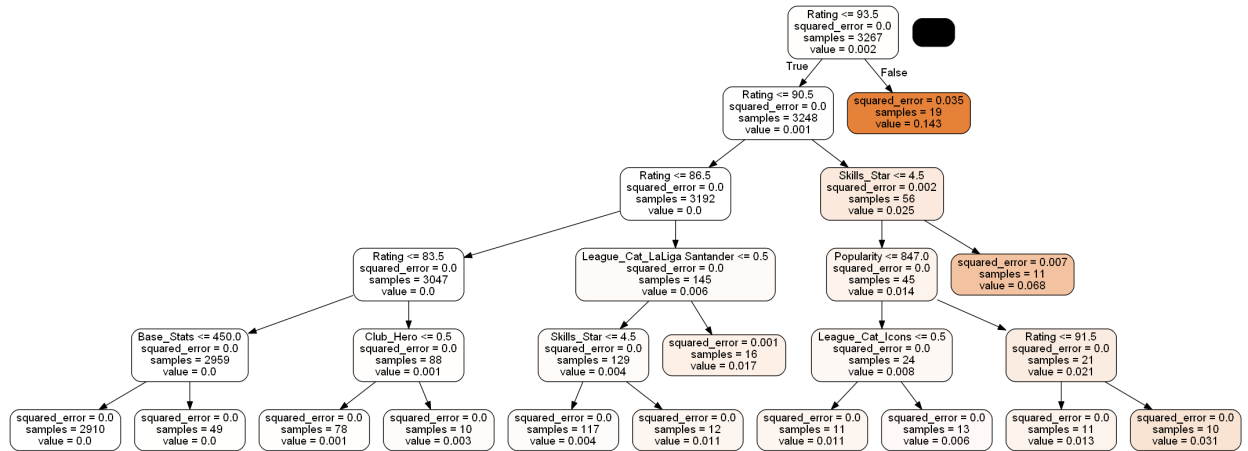
```
In [53]: rf.fit(X_train, y_train)
```

```
Out[53]: RandomForestRegressor(max_depth=5, min_samples_leaf=10, n_jobs=-1,
                             random_state=42)
```

```
In [54]: sample_tree = rf.estimators_[20]
```

```
In [55]: dot_data = StringIO()
export_graphviz(sample_tree, out_file=dot_data, filled=True, rounded=True,
               feature_names=X_train.columns)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

```
Out[55]:
```



```
In [56]: y_train_pred = rf.predict(X_train)
y_test_pred = rf.predict(X_test)
```

```
In [57]: r2_score(y_train, y_train_pred)
```

```
Out[57]: 0.4665518871340596
```

```
In [58]: r2_score(y_test, y_test_pred)
```

```
Out[58]: 0.6085673962637769
```

```
In [59]: rf.feature_importances_
```

```
Out[59]: array([8.63107840e-01, 4.11047310e-03, 0.00000000e+00, 2.35199905e-02,
        1.84573743e-04, 7.48589702e-03, 7.71977045e-02, 3.47173992e-04,
        4.02122200e-04, 1.96892230e-02, 3.62010511e-04, 9.44533176e-04,
        4.62558964e-06, 4.77114424e-04, 8.55871299e-04, 5.08042821e-05,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 4.59991441e-07, 3.69134191e-04, 5.66323717e-04,
        0.00000000e+00, 0.00000000e+00, 3.23781650e-07, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 2.36985890e-05, 2.04145436e-04, 2.39002353e-07,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 3.12153841e-05,
        0.00000000e+00, 0.00000000e+00, 5.30923721e-05, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 1.14105389e-05,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00])
```

```
In [60]: imp_df = pd.DataFrame({
        "Varname": X_train.columns,
        "Imp": rf.feature_importances_})
```

```
In [61]: imp_df.sort_values(by="Imp", ascending=False)
```

```
Out[61]:
```

	Varname	Imp
0	Rating	0.863108
6	Dribbling / Reflexes	0.077198
3	Pace / Diving	0.023520
9	Popularity	0.019689
5	Passing / Kicking	0.007486
...
43	BodyType_Text_Messi	0.000000
42	BodyType_Text_Lean	0.000000
41	BodyType_Text_High & Stocky	0.000000
40	BodyType_Text_High & Lean	0.000000
102	Alt_Pos_3_missing	0.000000

103 rows × 2 columns

Grid search for hyper-parameter tuning

```
In [62]: #from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [63]: classifier_rf = RandomForestRegressor(random_state=42, n_jobs=-1)
```

```
In [64]: # Create the parameter grid based on the results of random search
params = {
    'max_depth': [5, 8, 10, 13, 15, 17, 20],
    'min_samples_leaf': [5, 10, 20, 50, 100],
    'max_features': [2,3,4],
    'n_estimators': [10, 30, 50, 100, 200]
}
```

```
In [85]: # Instantiate the grid search model
# cv=4 -> 4-fold cross validation schema
grid_search = GridSearchCV(estimator=RandomForestRegressor(), param_grid=params,
                           cv=10, n_jobs=-1, verbose=1, scoring = "r2", error_score="raise")
```

```
In [70]: X_train.head()
```

```
Out[70]:
```

	Rating	Skills_Star	Weak_Foot_Star	Pace / Diving	Shooting / Handling	Passing / Kicking	Dribbling / Reflexes	Defense / Speed	Physical / Positioning
2556	69	2	3	68.0	64	65	68	66	71
5208	61	2	3	37.0	27	39	39	58	71
2567	69	2	2	68.0	37	63	66	67	68
3412	67	3	4	77.0	66	56	67	29	56
2673	69	2	3	73.0	67	65	70	36	55

5 rows × 103 columns

```
In [71]: y_train.head()
```

```
Out[71]: 2556    0.000043
5208    0.000653
2567    0.000113
3412    0.000027
2673    0.000017
Name: Price, dtype: float64
```

```
In [81]: X_train.shape
```

```
Out[81]: (5117, 103)
```

```
In [82]: y_train.shape
```

```
Out[82]: (5117,)
```

```
In [86]: %%time
grid_search.fit(X_train,y_train)
```

Fitting 10 folds for each of 525 candidates, totalling 5250 fits

CPU times: total: 25.4 s

Wall time: 7min 48s

```
Out[86]: GridSearchCV(cv=10, error_score='raise',
                    estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
                    n_jobs=-1,
                    param_grid={'max_depth': [5, 8, 10, 13, 15, 17, 20],
                                'max_features': [2, 3, 4],
                                'min_samples_leaf': [5, 10, 20, 50, 100],
                                'n_estimators': [10, 30, 50, 100, 200]},
                    scoring='r2', verbose=1)
```

```
In [87]: rf_best = grid_search.best_estimator_
```

```
In [88]: rf_best
```

```
Out[88]: RandomForestRegressor(max_depth=10, max_features=4, min_samples_leaf=5,
                               n_estimators=200, n_jobs=-1, random_state=42)
```

```
In [89]: y_train_pred = rf_best.predict(X_train)
         y_test_pred = rf_best.predict(X_test)
```

```
In [90]: r2_score(y_train, y_train_pred)
```

```
Out[90]: 0.42985622276084756
```

```
In [91]: r2_score(y_test, y_test_pred)
```

```
Out[91]: 0.43889540025568174
```

```
In [92]: rf_best.feature_importances_
```

```
Out[92]: array([1.27692642e-01, 4.83770707e-02, 1.59748454e-02, 7.20043965e-02,
                1.25897401e-01, 6.42622827e-02, 1.44817704e-01, 2.91795004e-02,
                2.34158031e-02, 5.72564589e-02, 6.34699221e-02, 6.74117801e-02,
                1.84441807e-02, 9.71723894e-03, 1.29137152e-02, 3.04718328e-03,
                1.73705341e-04, 2.03632739e-04, 8.11418527e-04, 1.51815146e-03,
                5.67718424e-05, 4.88507780e-05, 1.98656505e-05, 3.95510989e-03,
                5.93939023e-06, 3.87282426e-04, 1.68129795e-04, 6.87949271e-05,
                3.80398324e-06, 4.95817141e-03, 2.84880037e-03, 1.50803630e-03,
                3.05452165e-05, 2.59960910e-03, 8.65948284e-03, 2.93119221e-03,
                0.00000000e+00, 0.00000000e+00, 5.23205658e-04, 1.32026536e-04,
                0.00000000e+00, 6.17288287e-05, 7.75539661e-04, 0.00000000e+00,
                0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
                2.88387919e-06, 6.88842143e-05, 5.63710998e-07, 2.78378520e-05,
                6.65866942e-03, 2.54568151e-02, 4.19120803e-04, 5.38554725e-03,
                2.72782966e-05, 7.62377128e-03, 9.41663787e-04, 2.51730785e-04,
                3.51470124e-04, 5.85266245e-03, 3.27916229e-04, 5.31958718e-03,
                1.77003759e-04, 7.05235981e-06, 4.09450579e-04, 9.20927124e-03,
                1.01654689e-04, 1.66282089e-05, 1.76081732e-03, 1.45226799e-06,
                1.26045267e-05, 6.03409815e-05, 5.95764833e-04, 8.38360790e-06,
                8.82600803e-05, 3.80730845e-04, 1.41696119e-03, 4.68024100e-07,
                1.03565768e-09, 3.25459223e-05, 1.67554247e-05, 2.54097029e-08,
                3.58939524e-04, 2.80577548e-03, 3.45847481e-07, 0.00000000e+00,
                1.96593675e-05, 8.52184913e-04, 2.81891856e-04, 5.30810895e-03,
                2.24765404e-08, 1.99289790e-06, 0.00000000e+00, 0.00000000e+00,
                7.54046813e-07, 3.89393230e-04, 1.67210044e-07, 0.00000000e+00,
                1.19433210e-08, 2.59899766e-07, 6.35997383e-04])
```



```
In [93]: imp_df = pd.DataFrame({
    "Varname": X_train.columns,
    "Imp": rf_best.feature_importances_})
```

```
In [95]: imp_df.sort_values(by="Imp", ascending=False)
```

```
Out[95]:
```

	Varname	Imp
6	Dribbling / Reflexes	0.144818
0	Rating	0.127693
4	Shooting / Handling	0.125897
3	Pace / Diving	0.072004
11	Ingame_Stats	0.067412
...
87	Alt_Pos_2_RB	0.000000
99	Alt_Pos_3_RM	0.000000
37	BodyType_Text_Courtois	0.000000
36	BodyType_Text_CR7	0.000000
40	BodyType_Text_High & Lean	0.000000

103 rows × 2 columns

second grid search

- max depth testing around 10
- max feature around 4
- sample leaf around 5
- estimator around 200

```
In [176... # Create the parameter grid based on the results of random search
params = {
    'max_depth': [8, 9, 10, 11, 12, 13],
    'min_samples_leaf': [2,3,4,5,6,7,8],
    'max_features': [4,6,8,10],
    'n_estimators': [100, 150, 200, 250, 300]
}
```

```
In [177... # Instantiate the grid search model
# cv=4 -> 4-fold cross validation schema
grid_search = GridSearchCV(estimator=RandomForestRegressor(), param_grid=params,
                           cv=10, n_jobs=-1, verbose=1, scoring = "r2", error_score="raise")
```

```
In [98]: %%time
grid_search.fit(X_train,y_train)
```

Fitting 10 folds for each of 840 candidates, totalling 8400 fits
 CPU times: total: 39.3 s
 Wall time: 41min 22s

```
Out[98]: GridSearchCV(cv=10, error_score='raise',
                    estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
                    n_jobs=-1,
                    param_grid={'max_depth': [8, 9, 10, 11, 12, 13],
                                'max_features': [4, 6, 8, 10],
                                'min_samples_leaf': [2, 3, 4, 5, 6, 7, 8],
                                'n_estimators': [100, 150, 200, 250, 300]},
                    scoring='r2', verbose=1)
```

```
In [99]: rf_best = grid_search.best_estimator_
```

```
In [100... rf_best
```

```
Out[100]: RandomForestRegressor(max_depth=8, max_features=8, min_samples_leaf=2,
                                n_jobs=-1, random_state=42)
```

```
In [102... y_train_pred = rf_best.predict(X_train)
y_test_pred = rf_best.predict(X_test)
```

```
In [103... r2_score(y_train, y_train_pred)
```

```
Out[103]: 0.7268290034088154
```

```
In [104... r2_score(y_test, y_test_pred)
```

```
Out[104]: 0.6122313567377293
```

```
In [105... rf_best.feature_importances_
```

```
Out[105]: array([1.38489015e-01, 4.66858447e-02, 1.51138366e-02, 8.28239884e-02,
                  1.08141260e-01, 5.74164927e-02, 1.26643262e-01, 3.73862489e-02,
                  2.83327822e-02, 4.88247605e-02, 5.54882980e-02, 4.12670687e-02,
                  1.53507463e-02, 1.37700267e-02, 1.60956680e-02, 2.00895569e-03,
                  1.26215167e-04, 6.95812904e-06, 1.30401345e-03, 1.33142932e-03,
                  9.07435426e-05, 1.82934812e-04, 6.30403671e-05, 1.51820461e-02,
                  1.25060262e-06, 4.52118325e-05, 2.80795915e-04, 1.16865943e-04,
                  2.21889599e-07, 9.76897426e-03, 2.31602367e-03, 1.17541680e-04,
                  1.33593246e-05, 2.19611689e-03, 1.77063715e-02, 4.95317456e-03,
                  0.00000000e+00, 0.00000000e+00, 2.81857258e-04, 2.38660430e-05,
                  2.92761059e-07, 7.57767917e-09, 7.08016776e-04, 5.55441104e-03,
                  0.00000000e+00, 2.16383224e-02, 0.00000000e+00, 0.00000000e+00,
                  2.83678825e-06, 1.11221718e-05, 1.28650138e-07, 3.47949213e-06,
                  7.19721972e-03, 1.79522994e-02, 1.61142358e-04, 6.74598724e-03,
                  8.03629055e-05, 1.20973113e-03, 2.06139183e-03, 6.32214843e-05,
                  1.32951783e-04, 3.92989831e-03, 9.04149985e-06, 4.41874535e-03,
                  2.76849241e-04, 2.83567090e-06, 1.32693296e-03, 6.05009338e-03,
                  3.58926598e-04, 2.21020250e-05, 1.22392886e-02, 5.34552132e-10,
                  9.55474349e-05, 1.50988231e-04, 1.61348511e-03, 3.23384668e-05,
                  4.40615221e-05, 3.01687466e-04, 7.27018012e-04, 0.00000000e+00,
                  0.00000000e+00, 1.58692422e-07, 8.16173342e-06, 2.19814835e-06,
                  7.60430400e-04, 1.49256156e-03, 3.25608299e-07, 5.71065547e-09,
                  9.13209993e-09, 1.11797017e-04, 1.46286353e-04, 7.17550745e-03,
                  1.13647612e-09, 2.32284173e-05, 0.00000000e+00, 0.00000000e+00,
                  7.38528353e-06, 9.10773651e-04, 2.26690642e-09, 0.00000000e+00,
                  4.46735026e-06, 3.73632595e-07, 4.28629382e-03])
```

```
In [106... imp_df = pd.DataFrame({
    "Varname": X_train.columns,
    "Imp": rf_best.feature_importances_})
```

In [107... `imp_df.sort_values(by="Imp", ascending=False)`

Out[107]:

	Varname	Imp
0	Rating	0.138489
6	Dribbling / Reflexes	0.126643
4	Shooting / Handling	0.108141
3	Pace / Diving	0.082824
5	Passing / Kicking	0.057416
...
80	Alt_Pos_2_CDM	0.000000
95	Alt_Pos_3_LB	0.000000
44	BodyType_Text_R9	0.000000
46	BodyType_Text_Salah	0.000000
94	Alt_Pos_3_CM	0.000000

103 rows × 2 columns

third grid search

In [108... `# Create the parameter grid based on the results of random search`
`params = {`
 `'max_depth': [8, 9, 10],`
 `'min_samples_leaf': [1,2,3],`
 `'max_features': [7,8,9,10],`
 `'n_estimators': [100, 150, 200, 250, 300]`
`}`

In [110... `# Instantiate the grid search model`
`# cv=4 -> 4-fold cross validation schema`
`grid_search = GridSearchCV(estimator=RandomForestRegressor(), param_grid=params,`
 `cv=10, n_jobs=-1, verbose=2, scoring = "r2", error_score="raise")`

In [111... `%%time`
`grid_search.fit(X_train,y_train)`

Fitting 10 folds for each of 180 candidates, totalling 1800 fits
 CPU times: total: 8.41 s
 Wall time: 8min 35s

Out[111]: `GridSearchCV(cv=10, error_score='raise', estimator=RandomForestRegressor(),`
 `n_jobs=-1,`
 `param_grid={'max_depth': [8, 9, 10], 'max_features': [7, 8, 9, 10],`
 `'min_samples_leaf': [1, 2, 3],`
 `'n_estimators': [100, 150, 200, 250, 300]},`
 `scoring='r2', verbose=2)`

In [112... `rf_best = grid_search.best_estimator_`

In [113... rf_best

Out[113]: RandomForestRegressor(max_depth=9, max_features=10, min_samples_leaf=2)

In [114... y_train_pred = rf_best.predict(X_train)
y_test_pred = rf_best.predict(X_test)

In [115... r2_score(y_train, y_train_pred)

Out[115]: 0.755998374506422

In [116... r2_score(y_test, y_test_pred)

Out[116]: 0.6320821355740586

forth gridsearch

In [117... *# Create the parameter grid based on the results of random search*
params = {
 'max_depth': [8, 9, 10],
 'min_samples_leaf': [1, 2, 3],
 'max_features': [10, 20, 30, 40],
 'n_estimators': [100, 200, 300]
}

In [119... grid_search = GridSearchCV(estimator=RandomForestRegressor(), param_grid=params,
cv=10, n_jobs=-1, verbose=2, scoring = "r2", error_score="raise")

In [120... %%time
grid_search.fit(X_train, y_train)

Fitting 10 folds for each of 108 candidates, totalling 1080 fits

CPU times: total: 6.23 s

Wall time: 9min 41s

Out[120]: GridSearchCV(cv=10, error_score='raise', estimator=RandomForestRegressor(),
n_jobs=-1,
param_grid={'max_depth': [8, 9, 10],
'max_features': [10, 20, 30, 40],
'min_samples_leaf': [1, 2, 3],
'n_estimators': [100, 200, 300]},
scoring='r2', verbose=2)

In [121... rf_best = grid_search.best_estimator_

In [122... rf_best

Out[122]: RandomForestRegressor(max_depth=10, max_features=10, min_samples_leaf=2,
n_estimators=200)

In [123... y_train_pred = rf_best.predict(X_train)
y_test_pred = rf_best.predict(X_test)

In [124... r2_score(y_train, y_train_pred)

Out[124]: 0.7601409765037306

```
In [125... r2_score(y_test, y_test_pred)
```

```
Out[125]: 0.5793143634800904
```

fifth gridsearch

```
In [126... # Create the parameter grid based on the results of random search
params = {
    'max_depth': [8,9,10,11],
    'min_samples_leaf': [1,2,3],
    'max_features': [9,10,11,12],
    'n_estimators': [150,200,250]
}
```

```
In [127... grid_search = GridSearchCV(estimator=RandomForestRegressor(), param_grid=params,
                             cv=10, n_jobs=-1, verbose=2, scoring = "r2", error_score="raise")
```

```
In [128... %%time
grid_search.fit(X_train,y_train)
```

Fitting 10 folds for each of 144 candidates, totalling 1440 fits
CPU times: total: 6.61 s
Wall time: 7min 52s

```
Out[128]: GridSearchCV(cv=10, error_score='raise', estimator=RandomForestRegressor(),
                    n_jobs=-1,
                    param_grid={'max_depth': [8, 9, 10, 11],
                                'max_features': [9, 10, 11, 12],
                                'min_samples_leaf': [1, 2, 3],
                                'n_estimators': [150, 200, 250]},
                    scoring='r2', verbose=2)
```

```
In [129... rf_best = grid_search.best_estimator_
```

```
In [130... rf_best
```

```
Out[130]: RandomForestRegressor(max_depth=9, max_features=10, n_estimators=200)
```

```
In [131... y_train_pred = rf_best.predict(X_train)
y_test_pred = rf_best.predict(X_test)
```

```
In [132... r2_score(y_train, y_train_pred)
```

```
Out[132]: 0.9204138340558483
```

```
In [133... r2_score(y_test, y_test_pred)
```

```
Out[133]: 0.6727454133535952
```

Final model after grid search

```
In [166... rf = RandomForestRegressor(random_state=42, n_jobs=-1, max_depth=9, n_estimators=200,
```

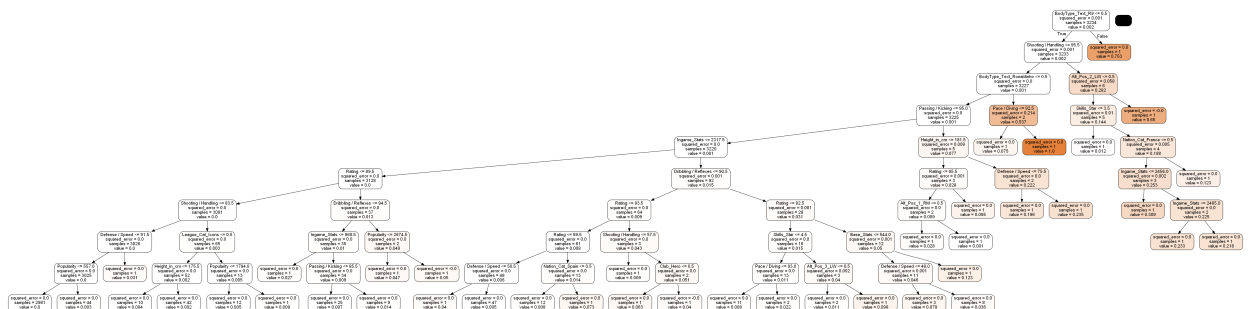
```
In [167... rf.fit(X_train, y_train)
```

```
Out[167]: RandomForestRegressor(max_depth=9, max_features=20, n_estimators=200, n_jobs=-1,
                             random_state=42)
```

```
In [168... sample_tree = rf.estimators_[2]
```

```
In [169... dot_data = StringIO()
export_graphviz(sample_tree, out_file=dot_data, filled=True, rounded=True,
                feature_names=X_train.columns)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

```
Out[169]:
```



```
In [170... y_train_pred = rf.predict(X_train)
y_test_pred = rf.predict(X_test)
```

```
In [171... r2_score(y_train, y_train_pred)
```

```
Out[171]: 0.9111099629582339
```

```
In [172... r2_score(y_test, y_test_pred)
```

```
Out[172]: 0.6946316867354081
```

```
In [173... rf.feature_importances_
```

Out[173]: array([1.54227577e-01, 3.38025014e-02, 6.13521060e-03, 8.57044644e-02, 9.32105450e-02, 3.49512604e-02, 1.28418526e-01, 3.77531647e-02, 1.37835014e-02, 4.18185863e-02, 4.12193749e-02, 3.46733427e-02, 1.68488482e-02, 1.44599500e-02, 8.03670500e-03, 1.09859057e-02, 1.10481212e-04, 9.16481082e-07, 3.02904448e-03, 7.34819817e-04, 8.84070374e-05, 2.85144951e-04, 1.36549923e-06, 1.26327298e-02, 1.53830061e-05, 1.41996130e-04, 9.62478940e-04, 1.93141729e-05, 3.23521504e-06, 3.55383930e-03, 4.86880557e-03, 7.02825543e-04, 1.25400345e-04, 2.12460059e-03, 8.77149995e-03, 3.19316754e-03, 5.27790193e-08, 0.00000000e+00, 1.48310230e-03, 1.81494113e-04, 6.16499733e-06, 2.43354677e-07, 7.88097204e-04, 7.13182824e-03, 3.55364412e-02, 7.36057306e-02, 1.13237740e-05, 0.00000000e+00, 5.86782212e-06, 2.68740766e-04, 5.62214915e-10, 3.32833972e-05, 6.32531925e-03, 9.40655932e-03, 7.35020932e-04, 6.00569782e-03, 1.59830095e-05, 8.61977285e-05, 2.41529934e-03, 6.86376700e-05, 9.53154197e-04, 3.76089087e-03, 6.04007599e-05, 2.38243433e-03, 4.12607711e-04, 4.41674367e-07, 9.74428991e-05, 6.25883367e-03, 2.44342043e-03, 7.07906689e-06, 1.42692791e-02, 9.48818672e-08, 1.10196878e-04, 2.09647352e-04, 2.49971385e-03, 7.49555006e-06, 1.18550056e-04, 2.11744875e-03, 2.42444882e-04, 9.44723831e-06, 0.00000000e+00, 1.20490703e-05, 7.77154166e-05, 2.57502802e-07, 2.18234774e-03, 7.69509295e-03, 4.07886068e-07, 9.14060277e-06, 1.57531629e-08, 1.68586791e-03, 4.10059193e-04, 2.67185498e-03, 3.07371070e-08, 7.52196774e-04, 0.00000000e+00, 0.00000000e+00, 7.65066083e-07, 2.55677591e-03, 0.00000000e+00, 0.00000000e+00, 6.07311500e-06, 1.49999476e-05, 4.45731962e-03])

```
In [174... imp_df = pd.DataFrame({
    "Varname": X_train.columns,
    "Imp": rf_best.feature_importances_})
```

```
In [175... imp_df.sort_values(by="Imp", ascending=False)
```

Out[175]:

	Varname	Imp
6	Dribbling / Reflexes	1.332667e-01
0	Rating	1.013360e-01
4	Shooting / Handling	9.981897e-02
3	Pace / Diving	6.632334e-02
45	BodyType_Text_Ronaldinho	6.168733e-02
...
47	BodyType_Text_Shaqiri	1.665979e-10
99	Alt_Pos_3_RM	1.368178e-13
94	Alt_Pos_3_CM	0.000000e+00
95	Alt_Pos_3_LB	0.000000e+00
80	Alt_Pos_2_CDM	0.000000e+00

103 rows × 2 columns

```
In [ ]:
```